

# POS tagging (3)

LING 570

Fei Xia

# Outline

- POS tagging with rich features
- Sequence labeling problem
- Beam search

# N-gram POS tagger

$$\operatorname{argmax}_{t_1^n} P(t_1^n | w_1^n)$$

$$\approx \operatorname{argmax}_{t_1^n} \prod_i P(w_i | t_i) P(t_i | t_{i-N+1}^{i-1})$$

Bigram model:  $\prod_i P(w_i | t_i) P(t_i | t_{i-1})$

Trigram model:  $\prod_i P(w_i | t_i) P(t_i | t_{i-2}, t_{i-1})$

# Unknown word handling

- HMM was good at using ***POS-tag context*** to pick POS for unknown words
- Bad at using information about the word itself
- Let's treat this as a classification problem:
  - Predict some target class
  - Use a bunch of ***features*** to make that prediction
  - ***Feature templates*** generate each feature

# POS Tagging with a classifier

- POS tagging as classification
  - What are the inputs?
    - What units are classified?
  - What are the classes?
  - What information should we use?

# Cues for unknown words

- Affixes: unforgettable: un-, -able → JJ
- Capitalization: Hyderabad → NNP
- Word shapes: 123,456 → CD
- The previous word: prevWord=San → NNP

How can we take advantage of these cues?

→ Treat them as features

# An example

- I am going to San Diego next week
- San NNP IsCap 1 PrevW=to 1 ContainNum 0
- Diego NNP IsCap 1 PrevW=San 1 ContainNum 0

# Feature templates for all the words

- Previous word:  $w_{-1}$
- Current word:  $w_0$
- Next word:  $w_{+1}$
- Previous two words:  $w_{-2} w_{-1}$
- Surrounding words:  $w_{-1} w_{+1}$
  
- Previous tag:  $t_{-1}$
- Previous two tags:  $t_{-2} t_{-1}$
  
- How many feature templates?
- How many features?  $3|V|+2|V|^2+|T|+|T|^2$



# An example

Mary will come tomorrow

	$W_{-1}$	$W_0$	$W_{-1} W_0$	$W_{+1}$	$t_{-1}$	$y$
x1 (Mary)	<s>	Mary	<s> Mary	will	BOS	PN
x2 (will)	Mary	will	Mary will	come	PN	V
x3 (come)	will	come	will come	tomorrow	V	V

This can be seen as a **shorthand** of a much bigger table.

	$W_{-1}$	$W_0$	$W_{-1} W_0$	$W_{+1}$	$t_{-1}$	$y$
x1 (Mary)	<s>	Mary	<s> Mary	will	BOS	PN
x2 (will)	Mary	will	Mary will	come	PN	V
x3 (come)	will	come	will come	tomorrow	V	V

Mary PN prevW=<s> 1 curW=Mary 1 prevW-curW=<s>-Mary 1  
nextW=will 1 prevTag=BOS 1

will V prevW=Mary 1 curW=will 1 prevW-curW=Mary-will 1  
nextW=come 1 prevTag=PN 1

come V prevW=will 1 curW=come 1 prevW-curW=will-come 1  
nextW=tomorrow 1 prevTag=V 1

# Ratnaparkhi's feature templates

Condition	Feature templates	
$w_i$ is not rare	$w_i = X$	$\& t_i = T$
$w_i$ is rare	$w_i$ has prefix $X$ , $ X  \leq 4$	$\& t_i = T$
	$w_i$ has suffix $X$ , $ X  \leq 4$	$\& t_i = T$
	$w_i$ contains number	$\& t_i = T$
	$w_i$ contains uppercase character	$\& t_i = T$
	$w_i$ contains hyphen	$\& t_i = T$
For all $w_i$	$t_{i-1} = X$	$\& t_i = T$
	$t_{i-2}, t_{i-1} = X, Y$	$\& t_i = T$
	$w_{i-1} = X$	$\& t_i = T$
	$w_{i-2} = X$	$\& t_i = T$
	$w_{i+1} = X$	$\& t_i = T$
	$w_{i+2} = X$	$\& t_i = T$

Word:	the	stories	about	well-heeled	communities	and	developers
Tag:	DT	NNS	IN	JJ	NNS	CC	NNS
Position:	1	2	3	4	5	6	7

Assume “well-heeled” is a rare word

```

well-heeled JJ  pref=w 1 pref=we 1 pref=wel 1 pref=well 1
                  suf=d 1 suf=ed 1 suf=led 1 suf=eled 1
                  containsNum 0 containsUppercase 0 containshyphen 1
                  prevTag=IN 1 prev2Tags=NNS-IN 1 prefW=about 1
                  pref2W=stories 1 nextW=communities 1 next2W=and 1

```

Rare words: words that occur less than  $N_r$  times in the training data

Feature selection: remove features that appear less than  $N_f$  times in the training data

# Building a tagger

- training data:  $w_1/t_1 \ w_2/t_2 \ \dots \ w_n/t_n$
- test data:  $w_1/t_1 \ w_2/t_2 \ \dots \ w_n/t_n$
- Steps:
  1. Create `train.vectors.txt` from training data
  2. Create `test.vectors.txt` from test data
  3. Run “mallet import-file” to convert training vectors to binary format
  4. Train a model using `train.vectors`:

```
mallet train-classifier --input train.vectors --trainer MaxEnt --output-classifier  
me_model --report train:accuracy > me.stdout 2>me.stderr
```
  5. Run the model on `test.vectors`:

```
mallet classify-file --input test.vectors.txt --classifier me_model --output  
resultFile --report test:accuracy > me_dec.stdout 2>me_dec.stderr
```
- Any problem?

	$W_{-1}$	$W_0$	$W_{-1} W_0$	$W_{+1}$	$t_{-1}$	$y$
x1 (Mary)	<s>	Mary	<s> Mary	will	BOS	PN
x2 (will)	Mary	will	Mary will	come	PN	V
x3 (come)	will	come	will come	tomorrow	V	V

Mary **PN** prevW=<s> 1 curW=Mary 1 prevW-curW=<s>-Mary 1  
 nextW=will 1 **prevTag=BOS** 1

will **V** prevW=Mary 1 curW=will 1 prevW-curW=Mary-will 1  
 nextW=come 1 **prevTag=PN** 1

come **V** prevW=will 1 curW=come 1 prevW-curW=will-come 1  
 nextW=tomorrow 1 **prevTag=V** 1

# Sequence labeling problem

# Sequence Labeling

- Classifier
  - Predict ***single output***, given potentially complex input
- Sequence classification
  - Predict ***sequence of output labels***, given sequence of potentially complex inputs



# Examples

- POS tagging
- NP chunking
- NE tagging
- Word segmentation
- Table detection
- ...

# Using a classifier

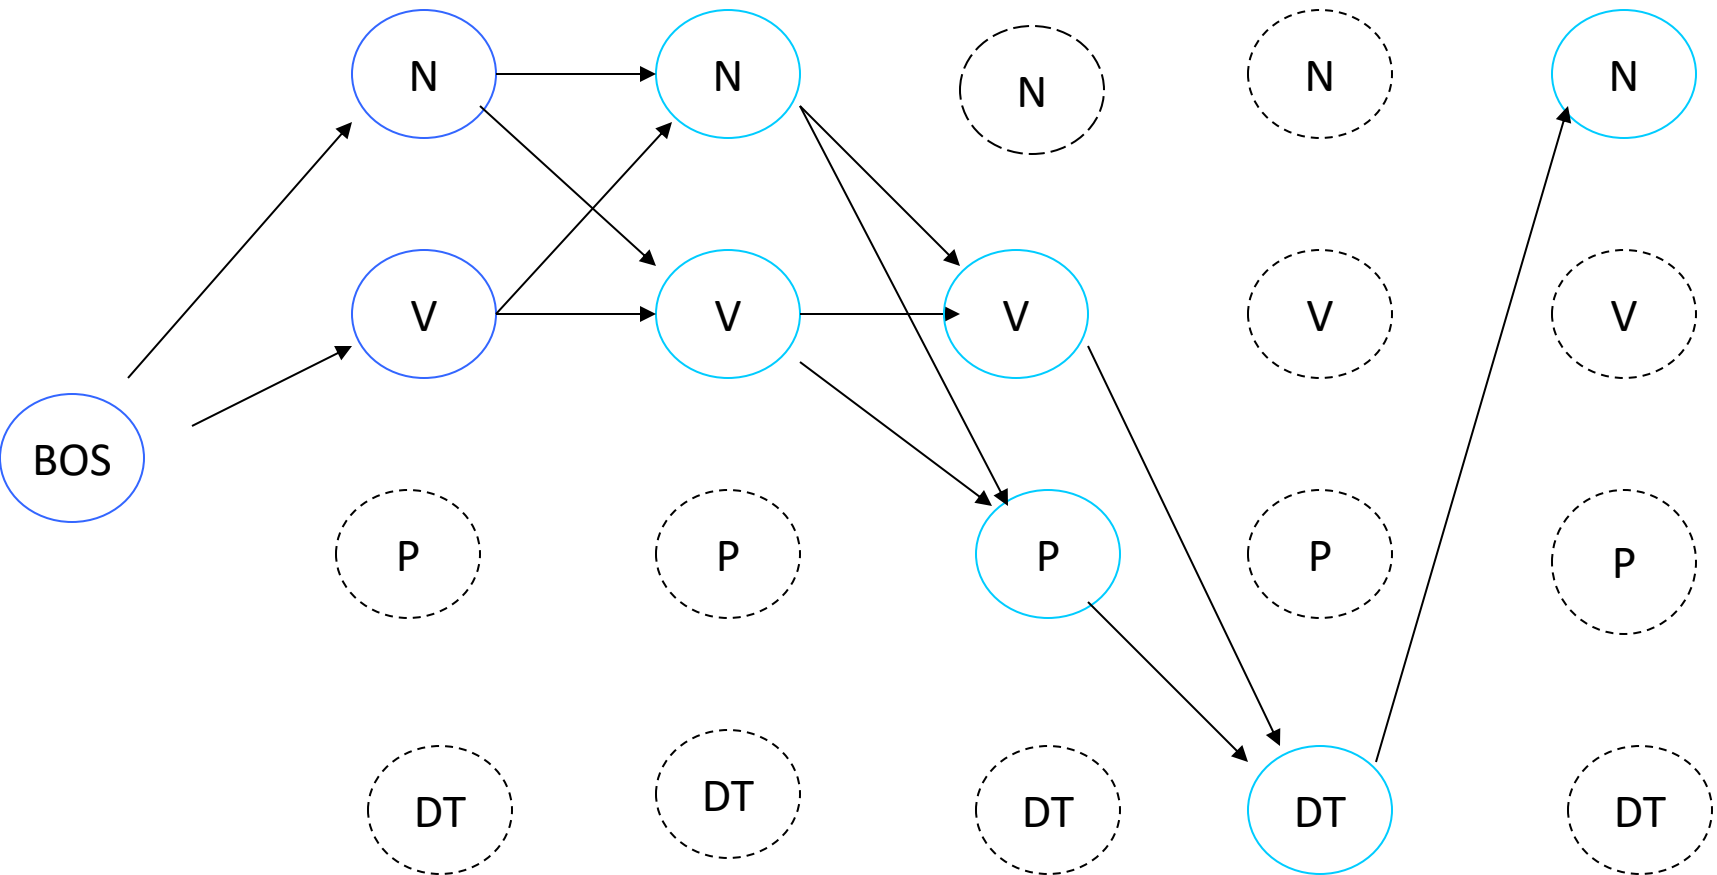
- Training data:  $\{(x_i, y_i)\}$
- What is  $x_i$ ? What is  $y_i$ ?
- What are the features?
- How to convert  $x_i$  to a feature vector for training data? How to do that for test data?


# How to solve a sequence labeling problem?

- Using a sequence labeling algorithm: e.g., HMM
- Using a classification algorithm:
  - Don't use features that refer to class labels
  - Use those features and get their values by running other processes
  - Use those features and find a good (global) solution.

# Viterbi for HMM

time flies like an arrow





$$\delta_j(t) \stackrel{\text{def}}{=} \max_{X_{1,t-1}} P(X_{1,t-1}, O_{1,t-1}, X_t = j)$$

$$\delta_j(1) = \pi_j$$

$$\delta_j(t+1) = \max_i \delta_i(t) a_{ij} b_{j o_t}$$

Time complexity:  $O(N^2 T)$

Can we use Viterbi for a classifier that uses tags of previous words?


# Beam search

# Why do we need beam search?

- Features refer to tags of previous words, which are not available for the TEST data.
- Knowing only the **best** tag of the previous word is not good enough.
- So let's keep multiple tag sequences available during **the decoding**.

# Beam search

time flies like an arrow





# Beam search

- Generate  $m$  tags for  $w_1$ , set  $s_{1j}$  accordingly
- For  $i=2$  to  $n$  ( $n$  is the sentence length)
  - Expanding: For each surviving sequence  $s_{(i-1),j}$ 
    - Generate  $m$  tags for  $w_i$ , given  $s_{(i-1),j}$  as previous tag context
    - Append each tag to  $s_{(i-1),j}$  to make a new sequence.
  - Pruning: keep only the top  $k$  sequences
- Return highest prob sequence  $s_{n1}$ .

# Beam search (basic)

- Beam inference:
  - At each position keep the top  $k$  complete sequences.
  - Extend each sequence in each local way.
  - The extensions compete for the  $k$  slots at the next position.
- Advantages:
  - Fast; and beam sizes of 3–5 are as good or almost as good as exact inference in many cases.
  - Easy to implement (no dynamic programming required).
- Disadvantage:
  - Inexact: the globally best sequence can fall off the beam.

# Viterbi search

- Viterbi inference:
  - Dynamic programming or memoization.
  - Requires small window of state influence (e.g., past two states are relevant).
- Advantage:
  - Exact: the global best sequence is returned.
- Disadvantage:
  - Harder to implement long-distance state-state interactions (but beam inference tends not to allow long-distance resurrection of sequences anyway).

# Viterbi vs. Beam search

- DP vs. heuristic search
- Global optimal vs. inexact
- Small window vs. big window for features

# Summary

- POS tagging with a classifier: use a classifier to determine the class of the word
- Sequence labeling problem: the feature of the current word depends on the tags of previous words
- Beam search: brute-force search with pruning